
Reproducible Training-Free Inference Acceleration on Pythia-70M

Ruoyu Fu
NLP Course Final Project
Code: GitHub repository

Abstract

Efficient inference is a practical bottleneck for autoregressive language models because attention computation and the KV cache grow with context length. This work studies training-free inference modifications on Pythia-70M: a post-hoc grouped-query attention (GQA) approximation and two KV cache compression strategies implemented with KVPress, Knorm and StreamingLLM. We evaluate exact teacher-forced perplexity, cached next-token perplexity, time to first token (TTFT), time per output token (TPOT), and throughput on WikiText-103 and a PG-19 sample. StreamingLLM-style KV compression improves WikiText throughput by 21.7–27.6% while preserving a reasonable cached-PPL score, whereas the post-hoc GQA approximation severely degrades perplexity and does not reliably accelerate inference. On PG-19, all compression variants slow down in our CPU-only runtime, showing that algorithmic cache reduction can be dominated by Python hook and gather overhead. The experiments can be reproduced with the scripts included in the repository.

1 Introduction

Autoregressive decoding repeatedly evaluates a transformer while maintaining a key-value (KV) cache for previous tokens. This cache avoids recomputing attention states, but its memory and attention cost scale with sequence length. Recent efficient-inference work therefore attacks the problem from several directions: using fewer KV heads (Ainslie et al., 2023), more IO-aware attention kernels (Dao et al., 2022), or compressing the KV cache during generation (Xiao et al., 2023; NVIDIA, 2024).

This project focuses on methods that can be reproduced without training a new model. We use Pythia-70M (Biderman et al., 2023), a small GPT-NeoX model, as the common backbone. The assignment asks for both quality and speed evaluation on datasets such as WikiText and PG-19, so we report two complementary quality metrics. First, exact perplexity is computed in the standard teacher-forced way and is not affected by KVPress, because KVPress acts on generation caches rather than full-context forward passes. Second, cached perplexity evaluates next-token prediction after a prefill pass whose KV cache may be compressed. This second metric is the appropriate quality signal for KVPress in this implementation.

The main contribution is a reproducible comparison under a CPU-only course environment. The results show that post-hoc GQA is not a reliable drop-in acceleration method for a model trained with standard multi-head attention, while KV cache compression can help on some contexts but is sensitive to runtime overhead.

2 Methods

Baseline. The baseline loads the local `EleutherAI/pythia-70m` checkpoint through Hugging Face Transformers and uses the default scaled dot-product attention (SDPA) implementation. No weights are changed.

Post-hoc GQA approximation. Pythia-70M uses 8 attention heads. In the GQA variant, we compute the original query, key, and value states, average the key/value heads into 2 grouped KV heads, then repeat the grouped KV states back to 8 heads before attention. This preserves tensor compatibility with the pretrained checkpoint, but it is only an inference-time approximation; the model was not trained under the grouped-KV constraint. We therefore expect possible quality degradation.

KVPress cache compression. We evaluate two KVPress modes at compression ratio 0.5. Knorm keeps tokens with larger key norm scores. StreamingLLM keeps attention-sink tokens and recent tokens, following the observation that early sink tokens are important for stable long-context decoding (Xiao et al., 2023). In both cases, compression is applied after the prefill stage via forward hooks on GPT-NeoX attention layers. The implementation records per-layer cache lengths before and after compression, e.g., 512 cached positions are compressed to 256.

Metrics. Exact PPL is computed over 4096 tokens using 512-token chunks. Cached PPL first prefills 512 context tokens, optionally compresses the resulting cache, and then scores the next 256 ground-truth tokens autoregressively with the cache. Speed is measured with greedy generation for 64 new tokens at context lengths 128, 512, and 1024. We report TTFT, TPOT, and throughput. All experiments in this report are single-run CPU measurements, so small differences should not be over-interpreted.

3 Experimental setup

We evaluate on the test split of WikiText-103 (Merity et al., 2016) and on one PG-19 test sample (Rae et al., 2020), matching the course requirement that PG-19 can be evaluated with a single long sample. The local environment is CPU-only with PyTorch 2.6.0, Transformers 4.56.2, Datasets 3.6.0, and KVPress 0.5.3. Since CUDA is unavailable, FlashAttention is not used as a main experimental condition.

The exact reproduction command is:

```
cd finalproj
.\scripts\run_matrix.ps1
```

This script runs baseline, GQA, KVPress-Knorm, and KVPress-StreamingLLM on both datasets and regenerates `comparison_quality.*` and `comparison_speed.*`.

4 Results and discussion

Table 1 summarizes quality. The exact-PPL column should be used to compare baseline and GQA. The cached-PPL column should be used to compare KVPress cache-compression behavior, because it scores continuation tokens after compressed prefill. GQA causes a large degradation: WikiText exact PPL rises from 63.72 to 2258.01, and PG-19 exact PPL rises from 33.57 to 816.69. This confirms that converting a pretrained MHA checkpoint into grouped KV heads by simple averaging is not quality-preserving.

Table 2 reports throughput relative to the baseline. On WikiText, KVPress-StreamingLLM gives the clearest acceleration: throughput improves by 26.7%, 21.7%, and 27.6% at context lengths 128, 512, and 1024, respectively. Knorm gives only small gains of 1.8–4.2%. GQA does not provide reliable speedup and also has poor quality, so it is mainly a negative result.

The PG-19 speed results are negative: every non-baseline method is slower. This does not mean KV compression is useless; rather, it shows that in this CPU implementation the overhead of Python

Table 1: Quality metrics. Exact PPL is standard teacher-forced PPL; cached PPL scores 256 continuation tokens after a 512-token prefill cache.

Dataset	Method	Exact PPL	Cached PPL	KV events	Avg. comp.
WikiText	Baseline	63.72	13.09	0	–
WikiText	GQA-2KV	2258.01	3658.94	0	–
WikiText	KVPress-Knorm	63.72	52.44	42	0.50
WikiText	KVPress-StreamingLLM	63.72	61.30	42	0.50
PG-19	Baseline	33.57	29.60	0	–
PG-19	GQA-2KV	816.69	776.88	0	–
PG-19	KVPress-Knorm	33.57	31.76	42	0.50
PG-19	KVPress-StreamingLLM	33.57	29.49	42	0.50

Table 2: Throughput change relative to baseline. Positive values indicate faster generation.

Dataset	Method	128 ctx	512 ctx	1024 ctx
WikiText	GQA-2KV	-2.2%	-14.0%	-7.5%
WikiText	KVPress-Knorm	+4.2%	+2.0%	+1.8%
WikiText	KVPress-StreamingLLM	+26.7%	+21.7%	+27.6%
PG-19	GQA-2KV	-58.2%	-86.7%	-88.9%
PG-19	KVPress-Knorm	-81.6%	-87.7%	-87.3%
PG-19	KVPress-StreamingLLM	-84.9%	-84.7%	-86.5%

hooks, tensor gathering, and cache manipulation can exceed the attention savings. This is especially plausible for Pythia-70M, where the model is small and the absolute attention workload is modest. A GPU implementation with fused kernels and larger models would be a more favorable setting for cache compression.

5 Conclusion

The most reproducible positive result is KVPress-StreamingLLM on WikiText, which improves throughput by about 22–28% while explicitly compressing each prefill cache by 50%. Knorm gives weaker speed gains. Post-hoc GQA is not successful: it sharply worsens perplexity and does not consistently accelerate generation. Overall, the study supports a cautious conclusion: training-free KV cache compression is a practical direction for inference acceleration, but its benefit depends on implementation overhead, dataset/runtime behavior, and model size. Negative results are important here because they prevent overstating a method that is theoretically appealing but not effective under the actual course environment.

References

- Biderman, S., Schoelkopf, H., Anthony, Q., Bradley, H., O’Brien, K., Hallahan, E., Khan, M. A., Purohit, S., Prashanth, U. S., Raff, E., Skowron, A., Sutawika, L., and Van Der Wal, O. Pythia: A suite for analyzing large language models across training and scaling. *International Conference on Machine Learning*, 2023.
- Ainslie, J., Lee-Thorp, J., de Jong, M., Zemlyanskiy, Y., Lebron, F., and Sanghai, S. GQA: Training generalized multi-query transformer models from multi-head checkpoints. *Empirical Methods in Natural Language Processing*, 2023.
- Dao, T., Fu, D. Y., Ermon, S., Rudra, A., and Re, C. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. *Advances in Neural Information Processing Systems*, 2022.
- Xiao, G., Tian, Y., Chen, B., Han, S., and Lewis, M. Efficient streaming language models with attention sinks. arXiv:2309.17453, 2023.
- NVIDIA. KVPress: KV cache compression for efficient LLM inference. <https://github.com/NVIDIA/kvpress>, 2024.
- Merity, S., Xiong, C., Bradbury, J., and Socher, R. Pointer sentinel mixture models. arXiv:1609.07843, 2016.
- Rae, J. W., Potapenko, A., Jayakumar, S. M., and Lillicrap, T. P. Compressive transformers for long-range sequence modelling. *International Conference on Learning Representations*, 2020.